

## Context

BankSim is an agent-based simulator of bank payments based on a sample of aggregated transactional data provided by a bank in Spain. The main purpose of BankSim is the generation of synthetic data that can be used for fraud detection research. Statistical and a Social Network Analysis (SNA) of relations between merchants and customers were used to develop and calibrate the model. Our ultimate goal is for BankSim to be usable to model relevant scenarios that combine normal payments and injected known fraud signatures. The data sets generated by BankSim contain no personal information or disclosure of legal and private customer transactions. Therefore, it can be shared by academia, and others, to develop and reason about fraud detection methods. Synthetic data has the added benefit of being easier to acquire, faster and at less cost, for experimentation even for those that have access to their own data. We argue that BankSim generates data that usefully approximates the relevant aspects of the real data.

## Content

We ran BankSim for 180 steps (approx. six months), several times and calibrated the parameters in order to obtain a distribution that get close enough to be reliable for testing. We collected several log files and selected the most accurate. We injected thieves that aim to steal an average of three cards per step and perform about two fraudulent transactions per day. We produced 594643 records in total. Where 587443 are normal payments and 7200 fraudulent transactions. Since this is a randomised simulation the values are of course not identical to original data.

```
In [12]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
In [13]: # Suppressing Warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [14]: # Importing Pandas and NumPy
import pandas as pd, numpy as np
```

```
In [15]: # Importing all datasets
fraud_detection = pd.read_csv("C:/Users/HP/Desktop/Fraud_Detection/fraud_detection.csv")
fraud_detection.head(4)
```

```
Out[15]:
```

	step	customer	age	gender	zipcodeOri	merchant	zipMerchant	category	amount	fraud
0	0	'C1093826151'	'4'	'M'	'28007'	'M348934600'	'28007'	'es_transportation'	4.55	0
1	0	'C352968107'	'2'	'M'	'28007'	'M348934600'	'28007'	'es_transportation'	39.68	0
2	0	'C2054744914'	'4'	'F'	'28007'	'M1823072687'	'28007'	'es_transportation'	26.89	0
3	0	'C1760612790'	'3'	'M'	'28007'	'M348934600'	'28007'	'es_transportation'	17.25	0

```
In [16]: fraud_detection.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 594643 entries, 0 to 594642
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   step             594643 non-null  int64  
1   customer         594643 non-null  object  
2   age              594643 non-null  object  
3   gender           594643 non-null  object  
4   zipcodeOri       594643 non-null  object  
5   merchant         594643 non-null  object  
6   zipMerchant      594643 non-null  object  
7   category         594643 non-null  object  
8   amount           594643 non-null  float64 
9   fraud            594643 non-null  int64  
dtypes: float64(1), int64(2), object(7)
memory usage: 45.4+ MB
```

```
In [17]: fraud_detection.shape
```

```
Out[17]: (594643, 10)
```

## FEATURE ENGINEERING

## Treating the Missing Values

Data That can be missing can be of two types :

- 1) Continuous Data
- 2) Discrete Or Categorical Data

The Types of missing can be of mentioned types :

- 1) **MCAR** - Missing Completely At Random

If the probability of being missing is same for all the observations.

- 2) **MNAR** - Missing Not At Random

There is some relationship between the missing data

- 3) **MAR** - Missing At Random

```
In [18]: fraud_detection.isnull().sum()
```

```
Out[18]: step          0
customer      0
age           0
gender        0
zipcodeOri    0
merchant      0
zipMerchant   0
category      0
amount        0
fraud         0
dtype: int64
```

There is no missing values for this data . In case if there is any missing values , we can use the mentioned below function for imputation

```
In [19]: def impute_nan(df,variable,median):
df[variable+"_median"]=df[variable].fillna(median)
df[variable+"_random"]=df[variable]
##It will have the random sample to fill the na
random_sample=df[variable].dropna().sample(df[variable].isnull().sum(),random_state=0)
##pandas need to have same index in order to merge the dataset
random_sample.index=df[df[variable].isnull()].index
df.loc[df[variable].isnull(),variable+'_'+random]=random_sample
```

```
In [20]: fraud_detection.head(4)
```

```
Out[20]:
```

	step	customer	age	gender	zipcodeOri	merchant	zipMerchant	category	amount	fraud
0	0	'C1093826151'	'4'	'M'	'28007'	'M348934600'	'28007'	'es_transportation'	4.55	0
1	0	'C352968107'	'2'	'M'	'28007'	'M348934600'	'28007'	'es_transportation'	39.68	0
2	0	'C2054744914'	'4'	'F'	'28007'	'M1823072687'	'28007'	'es_transportation'	26.89	0
3	0	'C1760612790'	'3'	'M'	'28007'	'M348934600'	'28007'	'es_transportation'	17.25	0

```
In [21]: fraud_detection.gender = fraud_detection.gender.astype('category').cat.codes
```

```
In [22]: fraud_detection.category = fraud_detection.category.astype('category').cat.codes
```

```
In [23]: fraud_detection.customer = fraud_detection.customer.astype('category').cat.codes  
fraud_detection.merchant = fraud_detection.merchant.astype('category').cat.codes
```

```
In [24]: fraud_detection.head(3)
```

```
Out[24]:
```

	step	customer	age	gender	zipcodeOri	merchant	zipMerchant	category	amount	fraud
0	0	210	'4'	2	'28007'	30	'28007'	12	4.55	0
1	0	2753	'2'	2	'28007'	30	'28007'	12	39.68	0
2	0	2285	'4'	1	'28007'	18	'28007'	12	26.89	0

```
In [25]: fraud_detection.age = fraud_detection.age.str.replace('[\\','], ', regex=True)
```

```
In [26]: fraud_detection.zipcodeOri = fraud_detection.zipcodeOri.str.replace('[\\','], ', regex=True)
```

```
In [27]: fraud_detection.zipMerchant = fraud_detection.zipMerchant.str.replace('[\\','], ', regex=True)
```

```
In [28]: fraud_detection.dtypes
```

```
Out[28]: step          int64
customer      int16
age           object
gender        int8
zipcodeOri    object
merchant      int8
zipMerchant   object
category      int8
amount        float64
fraud         int64
dtype: object
```

```
In [29]: fraud_detection["age"] = pd.to_numeric(fraud_detection["age"], errors='ignore')
```

```
In [30]: fraud_detection["zipcodeOri"] = pd.to_numeric(fraud_detection["zipcodeOri"], errors='coerce')
```

```
In [31]: fraud_detection["zipMerchant"] = pd.to_numeric(fraud_detection["zipMerchant"], errors='coerce')
```

```
In [32]: fraud_detection.head(3)
```

```
Out[32]:
```

	step	customer	age	gender	zipcodeOri	merchant	zipMerchant	category	amount	fraud
0	0	210	4	2	28007	30	28007	12	4.55	0
1	0	2753	2	2	28007	30	28007	12	39.68	0
2	0	2285	4	1	28007	18	28007	12	26.89	0

```
In [33]: fraud_detection.age.value_counts()
```

```
Out[33]: 2    187310
3    147131
4    109025
5     62642
1     58131
6     26774
0      2452
U       1178
Name: age, dtype: int64
```

HERE 'U' and '0' stands for error data. Drop the data

```
In [34]: fraud_detection = fraud_detection[~(fraud_detection.age == '0')]
```

```
In [35]: fraud_detection = fraud_detection[~(fraud_detection.age == 'U')]
```

```
In [36]: fraud_detection.age.value_counts()
```

```
Out[36]: 2    187310
         3    147131
         4    109025
         5     62642
         1     58131
         6     26774
         Name: age, dtype: int64
```

```
In [37]: X = fraud_detection.drop(columns="fraud")
```

```
In [38]: Y = fraud_detection["fraud"]
```

```
In [39]: fraud_detection.columns
```

```
Out[39]: Index(['step', 'customer', 'age', 'gender', 'zipcodeOri', 'merchant',
               'zipMerchant', 'category', 'amount', 'fraud'],
              dtype='object')
```


```
In [40]: fraud_detection.isnull().sum()
```

```
Out[40]: step          0
         customer      0
         age           0
         gender         0
         zipcodeOri     0
         merchant       0
         zipMerchant     0
         category       0
         amount         0
         fraud          0
         dtype: int64
```

```
In [41]: import matplotlib.pyplot as plt

plt.figure(figsize=(18,5))
plt.subplot(1,8,1)
plt.hist(fraud_detection['step'],bins=20)
plt.subplot(1,8,2)
plt.hist(fraud_detection['customer'],bins=20)
plt.subplot(1,8,3)
plt.hist(fraud_detection['age'],bins=20)
plt.subplot(1,8,4)
plt.hist(fraud_detection['gender'],bins=20)
plt.subplot(1,8,5)
plt.hist(fraud_detection['zipcodeOri'],bins=20)
plt.subplot(1,8,6)
plt.hist(fraud_detection['merchant'],bins=20)
plt.subplot(1,8,7)
plt.hist(fraud_detection['category'],bins=20)
plt.subplot(1,8,8)
plt.hist(fraud_detection['amount'],bins=20)

plt.show()
```



```
In [42]: # choosing all the numerical variables as independent variables (classifier can only take numerical input)
# dropping two variable funded_amnt as we have created new variable transformation based on it
X = fraud_detection.drop(columns = "fraud")
Y = fraud_detection["fraud"]

#splitting the dataset in train and test datasets using a split ratio of 70:30

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=10)
```

```
In [43]: print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)

(413709, 9) (413709,) (177304, 9) (177304,)
```

## FEATURE SELECTION

```
In [44]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
```

```
In [45]: from sklearn.ensemble import ExtraTreesClassifier
import matplotlib.pyplot as plt
model=ExtraTreesClassifier()
model.fit(X_train,y_train)
```


```
Out[45]: ▾ ExtraTreesClassifier
ExtraTreesClassifier()
```

```
In [46]: print(model.feature_importances_)

[0.07790335 0.07974806 0.03347501 0.00642493 0.          0.20321823
 0.          0.16733023 0.43190019]
```

```
In [47]: plt.figure(figsize = [10,4])
ranked_features=pd.Series(model.feature_importances_,index=X_train.columns)
ranked_features.nlargest(10).plot(kind='barh')
plt.show()
```





## Correlation - To Check Multicollinearity

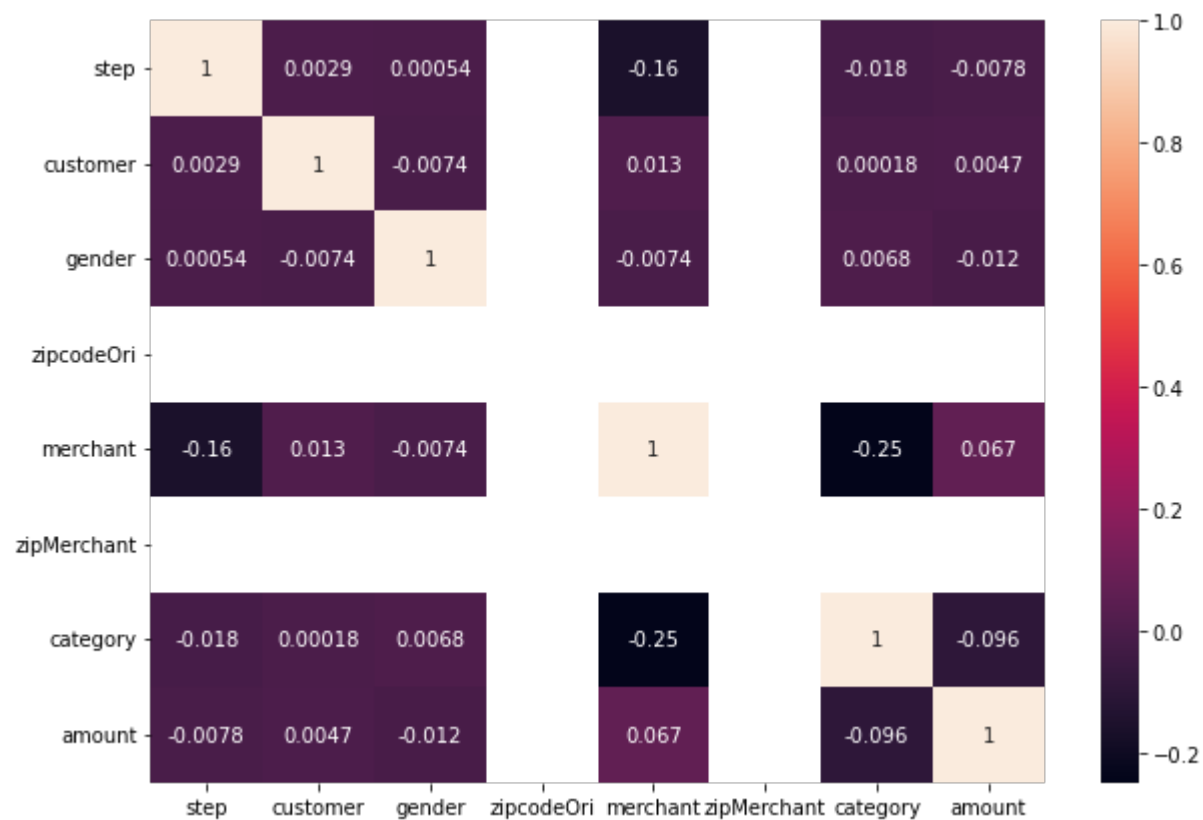
In [48]: `X_train.corr()`

Out[48]:

	step	customer	gender	zipcodeOri	merchant	zipMerchant	category	amount
step	1.000000	0.002905	0.000545	NaN	-0.155699	NaN	-0.018231	-0.007801
customer	0.002905	1.000000	-0.007449	NaN	0.013391	NaN	0.000175	0.004658
gender	0.000545	-0.007449	1.000000	NaN	-0.007365	NaN	0.006780	-0.012025
zipcodeOri	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
merchant	-0.155699	0.013391	-0.007365	NaN	1.000000	NaN	-0.247635	0.066526
zipMerchant	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
category	-0.018231	0.000175	0.006780	NaN	-0.247635	NaN	1.000000	-0.095956
amount	-0.007801	0.004658	-0.012025	NaN	0.066526	NaN	-0.095956	1.000000

In [49]: `import seaborn as sns  
corr=X_train.corr()  
top_features=corr.index  
plt.figure(figsize=(10,7))  
sns.heatmap(X_train[top_features].corr(),annot=True)`

Out[49]: <AxesSubplot:>



## Reduction Of Multi Collinearity

In [50]: threshold=0.6

```
In [51]: # find and remove correlated features
def correlation(dataset, threshold):
    col_corr = set() # Set of all the names of correlated columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i, j]) > threshold: # we are interested in absolute coeff value
                colname = corr_matrix.columns[i] # getting the name of column
                col_corr.add(colname)
    return col_corr
```

```
In [52]: correlation(X_train, threshold)
```

```
Out[52]: set()
```

```
In [53]: X.columns
```

```
Out[53]: Index(['step', 'customer', 'age', 'gender', 'zipcodeOri', 'merchant',
               'zipMerchant', 'category', 'amount'],
              dtype='object')
```

```
In [54]: X.drop(columns="zipcodeOri", inplace=True)
```

```
In [55]: X.drop(columns="zipMerchant", inplace=True)
```

```
In [56]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=10)
```

```
In [57]: print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)

(413709, 7) (413709,) (177304, 7) (177304,)
```

## PIPELINE CREATION

```
In [58]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
In [59]: ###! pip install xgboost
```

```
In [60]: from sklearn.preprocessing import StandardScaler,MinMaxScaler,MaxAbsScaler,RobustScaler
```

```
In [61]: pipeline_randomforest=Pipeline([('scalar3',RobustScaler()),
                                          ('pca3',PCA(n_components=2)),
                                          ('rf_classifier',RandomForestClassifier())])
```

```
In [62]: pipeline_gradient_boost=Pipeline([('scalar4',RobustScaler()),
                                             ('pca4',PCA(n_components=2)),
                                             ('gb_classifier',GradientBoostingClassifier())])
```

```
In [63]: pipeline_XGboost=Pipeline([('scalar5',RobustScaler()),
                                     ('pca5',PCA(n_components=2)),
                                     ('xgb_classifier',XGBClassifier())])
```

```
In [64]: ## Lets make the list of pipelines
pipelines = [pipeline_randomforest,pipeline_gradient_boost,pipeline_XGboost]
```

```
In [65]: best_accuracy=0.0
best_classifier=0
best_pipeline=""
```

```
In [66]: # Dictionary of pipelines and classifier types for ease of reference
pipe_dict = {0: 'RandomForest', 1: 'Gradient Boost', 2: 'Extreme Gradient Boost'}

# Fit the pipelines
for pipe in pipelines:
    pipe.fit(X_train, y_train)
```

```
[13:57:54] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```
In [67]: for i,model in enumerate(pipelines):  
         print("{} Test Accuracy: {}".format(pipe_dict[i],model.score(X_test,y_test)))
```

```
RandomForest Test Accuracy: 0.9952341740739069  
Gradient Boost Test Accuracy: 0.9948562920182286  
Extreme Gradient Boost Test Accuracy: 0.9955612958534494
```

```
In [68]: y_test.value_counts()
```

```
Out[68]: 0    175227  
         1     2077  
         Name: fraud, dtype: int64
```

```
In [69]: gb = GradientBoostingClassifier(n_estimators=100)  
         gb.fit(X_train, y_train)  
         preds = gb.predict(X_test)
```

```
In [70]: from sklearn import metrics
```

```
In [71]: # Confusion matrix  
         confusion = metrics.confusion_matrix(y_test, preds)  
         print(confusion)
```

```
[[175060   167]  
 [   578  1499]]
```

```
In [72]: random_grid = {'bootstrap': [True, False],  
                        'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None],  
                        'max_features': ['auto', 'sqrt'],  
                        'min_samples_leaf': [1, 2, 4],  
                        'min_samples_split': [2, 5, 10],  
                        'n_estimators': [130, 180, 230]}
```

```
In [73]: from sklearn.model_selection import RandomizedSearchCV
```

```
In [74]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator = rf,
                               param_distributions = random_grid,
                               n_iter = 5,
                               cv = 4,
                               verbose=2,
                               random_state=42,
                               n_jobs = -1)
```

```
In [75]: # Fit the random search model
rf_random.fit(X_train, y_train)
```

Fitting 4 folds for each of 5 candidates, totalling 20 fits

```
Out[75]: RandomizedSearchCV
          estimator: RandomForestClassifier
              RandomForestClassifier
```

```
In [76]: rf_random.best_params_
```

```
Out[76]: {'n_estimators': 130,
          'min_samples_split': 10,
          'min_samples_leaf': 2,
          'max_features': 'auto',
          'max_depth': 90,
          'bootstrap': False}
```

```
In [77]: best_random_grid=rf_random.best_estimator_
```

```
In [78]: from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
y_pred=best_random_grid.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print("Accuracy Score {}".format(accuracy_score(y_test,y_pred)))
print("Classification report: {}".format(classification_report(y_test,y_pred)))
```

```
[[175057 170]
 [ 537 1540]]
Accuracy Score 0.9960124983079908
Classification report:

```

			precision	recall	f1-score	support
	0	1.00	1.00	1.00		175227
	1	0.90	0.74	0.81		2077
accuracy				1.00		177304
macro avg		0.95	0.87	0.91		177304
weighted avg		1.00	1.00	1.00		177304

In [ ]:

In [ ]: